

Canon — making marketing claims provably true

A short brief · for review

A working system, not a concept: this is running today inside one of our own ventures (a regulated services platform) and is being extracted as a reusable tool any company can adopt. The example below uses real, public marketing copy to show how it would work for an adventure-travel brand.

Executive summary

Every growing company hits the same fight: the marketing person writes copy that sounds great, the founder isn't sure it's all true, and nobody can say *exactly* which claims are safe. So you either ship overpromises (refund disputes, "we handle visas" you don't, an expired certification badge still on the homepage) or you slow everything down arguing line by line.

Canon turns that argument into a fact-check a computer runs automatically. You write down three short lists once: what you've *actually* built and can deliver, what you're *allowed* to say, and what you must *never* say. Every claim you publish must trace to something real on those lists — or it simply doesn't go live. The website build fails and tells you which claim isn't backed, and why.

The effect: marketing keeps its speed and creativity, the founder stops being the human compliance filter, and the company can't quietly drift into claims it can't honor — even as the team grows and nobody remembers why a line was risky. It's the discipline of "don't say what isn't true" made mechanical, so it survives scale.

The problem (the "marketing guy" problem)

It's not that the marketer is careless. It's that **nobody owns a single source of truth about what's actually true**. So claims get made from optimism, get copied across pages, and rot quietly:

- A *guarantee* gets written that operations won't actually honor → refund disputes, chargebacks, bad reviews.
- A *capability* gets claimed ("we handle your visa") that's really "we refer you to a partner" → angry customers, possibly a regulatory issue.
- A *statistic* ("99.6% rate us 5 stars") goes up with no source on file → indefensible if challenged.
- A *certification badge* stays on the site after it lapses → a real liability nobody noticed.

Each is small. Together they're the difference between a brand people trust and one that gets caught out.

What it is (in plain terms)

Three short files you keep up to date — no code to read:

1. **What's real** — the honest list of what you've built and can actually deliver today (and what's only "coming soon").
2. **What you're allowed to say** — your approved claims. Each one is *linked* to an item on list 1. If the real thing isn't there, the claim isn't allowed.
3. **What you must never say** — banned phrases (e.g. "guaranteed," "risk-free," unverifiable superlatives) with the safe wording to use instead.

Plus, for a travel business, a fourth: **what you may sell, where** — because things like visas and insurance are regulated, and "we arrange it" vs "we handle it" is a legal line, not a style choice.

The one rule the machine enforces: *a claim goes public only if (a) it maps to something real you can deliver, and (b) it isn't a banned phrase.* Break either and the publish is blocked, with the exact reason.

Worked example — an adventure-travel brand (using live homepage copy)

The "blocked" verdicts are illustrative — each clears the moment the real capability or evidence is on file. That's the whole point.

Claim on the site	Verdict	Why
"100% Best Price Guaranteed"	Blocked	"Guaranteed" + a price promise must map to a real, written price-match process ops will honor. No documented process on file = no claim. Fix: confirm the process, or reword to "We'll match a lower price — ask us."
"100% refund guaranteed, first 24h, no questions asked"	Blocked until confirmed	A hard refund promise must trace to a refund policy finance actually executes. If "no questions asked" isn't literally how it works, this is a dispute waiting to happen. Confirm the policy → claim clears.
"We handle everything — flights, visas, transfers, insurance"	Blocked (regulated)	"Handle visas / insurance" claims authority you may not hold. If it's really "we help arrange via a partner," the gate forces the honest wording. This is the one that becomes a legal problem at scale.
"99.6% of our customers rate us 5 stars"	Needs a source	A specific number must have evidence attached (where it's from, when measured). Unsourced = flagged. Add the source → it ships, and now it's defensible.
"ISO 14001 & ISO 9001 certified"	Auto-expires	Certifications carry an expiry date. The gate warns and then blocks when the date passes, so an expired badge can never sit silently on the homepage.
"24/7 assistance + AI support"	Blocked if not live	If AI support is actually just WhatsApp today, "AI support" is marked <i>coming soon</i> and can't be claimed as shipped — until it is.
"Top-rated adventures travelers love" / "travel with the best"	Allowed	Brand voice, not a measurable promise. The gate distinguishes flavor from factual claims and lets the flavor through.

The valuable part isn't the red marks — it's that **the decision is no longer an opinion**. The founder and the marketer stop arguing about whether a line is OK. They look at one file. If the capability is real, the claim ships. If not, the fix is obvious: make it real, add the evidence, or reword. The fight becomes a checklist.

How it runs, day to day

1. Marketing writes whatever they want — full creative freedom.
2. Before it goes live, an automatic check confirms every claim traces to something real and avoids banned phrases.

3. Anything unbacked fails the publish with a plain-English reason. Make it real, add a source, or reword — then it ships.

No new tool to log into; it runs in the background of the normal publish process. The three lists take an afternoon to write the first time and minutes to maintain.

What I'd love your feedback on

- Does the "**argument becomes a checklist**" framing match your real pain — or is your problem more about the *quality / consistency* of the copy than its *truthfulness*?
- Would the **expiring-certification** and **unsourced-statistic** catches alone be worth it to you?
- Are you (the founder) willing to maintain the "what's actually true" list — or does that itself need to be someone's job?
- Would you want this as something **you run**, or as a **done-for-you service** — we set up the three lists with you, then it self-checks from then on?